

Stat Computing - Exercises 02 – Gradebook

Nick Gembs – ng334

There is a dataset in `datasets/grade_book.csv` containing simulated grades for Statistical Computing. From the exercises directory, you can read in the dataset by running

```
dat <- read.csv("../grade_book.csv")
str(dat)

## 'data.frame':    50 obs. of  26 variables:
## $ lab08 : int  8 8 6 6 8 5 10 8 8 8 ...
## $ lab03 : int  NA 5 6 4 7 5 7 8 8 8 ...
## $ lab01 : int  10 9 7 9 5 9 8 9 9 7 ...
## $ ex08 : int  NA 19 19 18 17 16 19 19 20 20 ...
## $ lab09 : int  9 9 8 9 7 7 8 4 9 8 ...
## $ ex09 : int  18 20 18 20 NA 19 19 20 18 18 ...
## $ ex04 : int  17 14 13 10 10 11 17 14 13 16 ...
## $ lab06 : int  7 9 7 8 6 10 9 7 7 8 ...
## $ project: int  82 83 NA 71 71 59 87 80 70 84 ...
## $ ex07 : int  19 18 17 19 16 12 17 17 14 19 ...
## $ ex03 : int  16 11 11 13 13 14 15 15 13 14 ...
## $ lab07 : int  9 7 6 8 10 6 8 9 6 10 ...
## $ ex12 : int  19 18 16 19 19 16 18 16 19 20 ...
## $ netID : chr  "gnz598" "hpl316" "ivp353" "iau101" ...
## $ lab10 : int  9 8 6 7 4 5 10 9 6 9 ...
## $ ex10 : int  15 16 NA 15 10 14 15 18 15 17 ...
## $ lab12 : int  9 8 6 9 7 6 9 7 9 7 ...
## $ ex02 : int  15 17 NA 15 15 16 15 19 17 19 ...
## $ ex05 : int  17 19 NA NA NA 17 18 17 13 17 ...
## $ lab05 : int  5 6 4 4 4 6 9 4 4 6 ...
## $ ex11 : int  16 15 15 18 18 18 19 NA 17 18 ...
## $ ex01 : int  12 15 11 16 13 12 14 16 11 11 ...
## $ ex06 : int  19 18 13 18 18 15 15 19 19 20 ...
## $ lab11 : int  9 9 8 8 10 7 10 10 8 8 ...
## $ lab04 : int  8 7 7 9 8 7 8 8 9 10 ...
## $ lab02 : int  8 9 6 10 10 10 9 10 10 10 ...
```

Complete the following exercises related to the grade book.

1. Randomly generate with replacement a birth date for each student between 2001-01-01 and 2005-12-31. Print out how many unique birth dates there are (year-month-date) and how many unique birthdays there are (month-date). Is this surprising?

```
dim(dat)

## [1] 50 26

bday = (as.Date(runif(50, min = 0, max = 1825), origin = as.Date("2001-01-01")))
```

```
bdayshort = format(bday, format = "%m/%d")
length(unique(bday))
```

```
## [1] 50
```

```
b = length(unique(bdayshort))
b
```

```
## [1] 46
```

There are 50 unique (year-month-date) birthdays and 46 unique (month-date) birthdays. This makes sense because there are 1825 unique birthday options with year while only 365 without year.

2. Add the birthdate column to the data frame in 3 different ways: using cbind, \$, and [[]].

```
way1 = cbind(dat, birthdate = bday)
```

```
#way2
```

```
dat$birthdate = bday
```

```
#way3
```

```
dat[["birthdate"]] = bday
```

3. Print out the column names.

```
names(dat)
```

```
## [1] "lab08" "lab03" "lab01" "ex08" "lab09" "ex09"
## [7] "ex04" "lab06" "project" "ex07" "ex03" "lab07"
## [13] "ex12" "netID" "lab10" "ex10" "lab12" "ex02"
## [19] "ex05" "lab05" "ex11" "ex01" "ex06" "lab11"
## [25] "lab04" "lab02" "birthdate"
```

4. Remove the birthdate column and show that it's gone.

```
dat$birthdate = NULL
```

```
names(dat)
```

```
## [1] "lab08" "lab03" "lab01" "ex08" "lab09" "ex09" "ex04"
## [8] "lab06" "project" "ex07" "ex03" "lab07" "ex12" "netID"
## [15] "lab10" "ex10" "lab12" "ex02" "ex05" "lab05" "ex11"
## [22] "ex01" "ex06" "lab11" "lab04" "lab02"
```

5. Print out which column number has the netIDs.

```
which( names(dat)=="netID" )
```

```
## [1] 14
```

6. Print out the column numbers that contain lab grades. You might find the grep function useful.

```
grep('lab', names(dat))
```

```
## [1] 1 2 3 5 8 12 15 17 20 24 25 26
```

7. Print out the column numbers that contain exercise grades.

```
grep('ex', names(dat))
```

```
## [1] 4 6 7 10 11 13 16 18 19 21 22 23
```

8. What happens when you try to convert the data frame to a matrix with `as.matrix`?

```
matdat <- as.matrix(dat)
```

It automatically converts all of the data types to string because a matrix has to have uniform data type.

9. Extract the exercise columns and convert to a matrix. Why does this work as intended?

```
datmat = as.matrix(dat[grep('ex', names(dat))])
```

This works as intended because the rule of uniform data type of a matrix is not broken, so R does not auto convert.

10. Add a column to the data frame containing each student's average exercise grade. Treat missing values as a grade of 0. You can do this in a couple of lines with `rowSums` or `rowMeans`. Exercises are out of 20. Print out the average exercise grades for the first 10 students.

```
datmat <- replace(datmat, is.na(datmat), 0)
dat$exAVG = 5*rowMeans(datmat, na.rm = TRUE)
5*rowMeans(datmat)[1:10]
```

```
## [1] 76.25000 83.33333 55.41667 75.41667 62.08333 75.00000 83.75000
79.16667
## [9] 78.75000 87.08333
```

11. Calculate each student's exercise average again, this time using the average of the non-missing values. Print out the average exercise grades for the first 10 students.

```
datmat = as.matrix(dat[grep('ex', names(dat))])
5*rowMeans(datmat, na.rm = TRUE)[1:10]
```

```
## [1] 108.02083 108.97436 94.20833 106.84028 95.94697 98.07692 109.51923
## [8] 112.15278 102.98077 113.87821
```

12. Print out the number of missing exercises for each exercise.

```
sum(is.na(datmat[,1]))
```

```
## [1] 1
```

```
sum(is.na(datmat[,2]))
```

```
## [1] 3
```

```
sum(is.na(datmat[,3]))
```

```
## [1] 1
```

```

sum(is.na(datmat[,4]))
## [1] 2
sum(is.na(datmat[,5]))
## [1] 1
sum(is.na(datmat[,6]))
## [1] 2
sum(is.na(datmat[,7]))
## [1] 1
sum(is.na(datmat[,8]))
## [1] 1
sum(is.na(datmat[,9]))
## [1] 5
sum(is.na(datmat[,10]))
## [1] 4

```

13. Calculate each student's lab average, and add to the data frame. Labs are out of 10. Print out the average lab grades for the first 10 students.

```

datmat = as.matrix(dat[grep('lab', names(dat))])

10*rowMeans(datmat, na.rm =TRUE)[1:10]

## [1] 82.72727 78.33333 64.16667 75.83333 71.66667 69.16667 87.50000
## [9] 77.50000 82.50000

#datmat <- replace(datmat, is.na(datmat), 0)
dat$labAVG = 10*rowMeans(datmat, na.rm =TRUE)

```

14. Using the formula in the syllabus, add a column containing each student's overall numeric grade. Treat missing assignments as 0. Project is out of 100.

```

dat <- replace(dat, is.na(dat), 0)

dat$numericgrade = dat$labAVG*.2 + dat$exAVG*.6 + dat$project*.2

```

15. Using the guidelines in the syllabus, add a column containing each student's letter grade.

```
lam = c()
for(i in 1:50){
  if(dat[i,"numericgrade"] >= 93){
    lam[i] = "A"
  } else if (dat[i,"numericgrade"] >= 90){
    lam[i] = "A-"
  } else if (dat[i,"numericgrade"] >= 87){
    lam[i] = "B+"
  } else if (dat[i,"numericgrade"] >= 83){
    lam[i] = "B"
  } else if (dat[i,"numericgrade"] >= 80){
    lam[i] = "B-"
  } else if (dat[i,"numericgrade"] >= 77){
    lam[i] = "C+"
  } else if (dat[i,"numericgrade"] >= 73){
    lam[i] = "C"
  } else if (dat[i,"numericgrade"] >= 70){
    lam[i] = "C-"
  } else if (dat[i,"numericgrade"] >= 67){
    lam[i] = "D+"
  } else if (dat[i,"numericgrade"] >= 63){
    lam[i] = "D"
  } else if (dat[i,"numericgrade"] >= 60){
    lam[i] = "D-"
  } else {
    lam[i] = "F"
  }
}

dat$lettergrade = lam

head(dat)
```

	lab08	lab03	lab01	ex08	lab09	ex09	ex04	lab06	project	ex07	ex03	lab07
## 1	8	0	10	0	9	18	17	7	82	19	16	9
## 2	8	5	9	19	9	20	14	9	83	18	11	7
## 3	6	6	7	19	8	18	13	7	0	17	11	6
## 4	6	4	9	18	9	20	10	8	71	19	13	8
## 5	8	7	5	17	7	0	10	6	71	16	13	10

```
## 6      5      5      9      16      7      19      11      10      59      12      14      6
16
##      netID lab10 ex10 lab12 ex02 ex05 lab05 ex11 ex01 ex06 lab11 lab04 lab02
## 1 gnz598      9      15      9      15      17      5      16      12      19      9      8      8
## 2 hpl316      8      16      8      17      19      6      15      15      18      9      7      9
## 3 ivp353      6      0      6      0      0      4      15      11      13      8      7      6
## 4 iau101      7      15      9      15      0      4      18      16      18      8      9      10
## 5 nue991      4      10      7      15      0      4      18      13      18      10      8      10
## 6 yky774      5      14      6      16      17      6      18      12      15      7      7      10
##      exAVG      labAVG      numericgrade      lettergrade
## 1 76.25000 82.72727      78.69545      C+
## 2 83.33333 78.33333      82.26667      B-
## 3 55.41667 64.16667      46.08333      F
## 4 75.41667 75.83333      74.61667      C
## 5 62.08333 71.66667      65.78333      D
## 6 75.00000 69.16667      70.63333      C-
```

16. Print out the netID, numeric average, and letter grade for the top 10 scorers. You may want to look at the order function.

```
tops = dat[order(dat$numericgrade,
decreasing=TRUE),c("netID", "numericgrade", "lettergrade")]
tops[1:10,]
```

```
##      netID numericgrade      lettergrade
## 22 esy224      92.26667      A-
## 24 waq733      89.78333      B+
## 36 hbz284      89.76667      B+
## 48 ayp949      89.35000      B+
## 42 bor334      86.86667      B
## 21 ujq876      86.61667      B
## 34 rtq675      86.10000      B
## 43 sxz212      85.83333      B
## 10 kap440      85.55000      B
## 7  pmc842      85.15000      B
```